Alpha runs on a laptop; only training needs a GPU

Level 01 - local deployment. Nothing here needs an AWS account, which keeps the $200 credit for the Beta stack (risk R17).

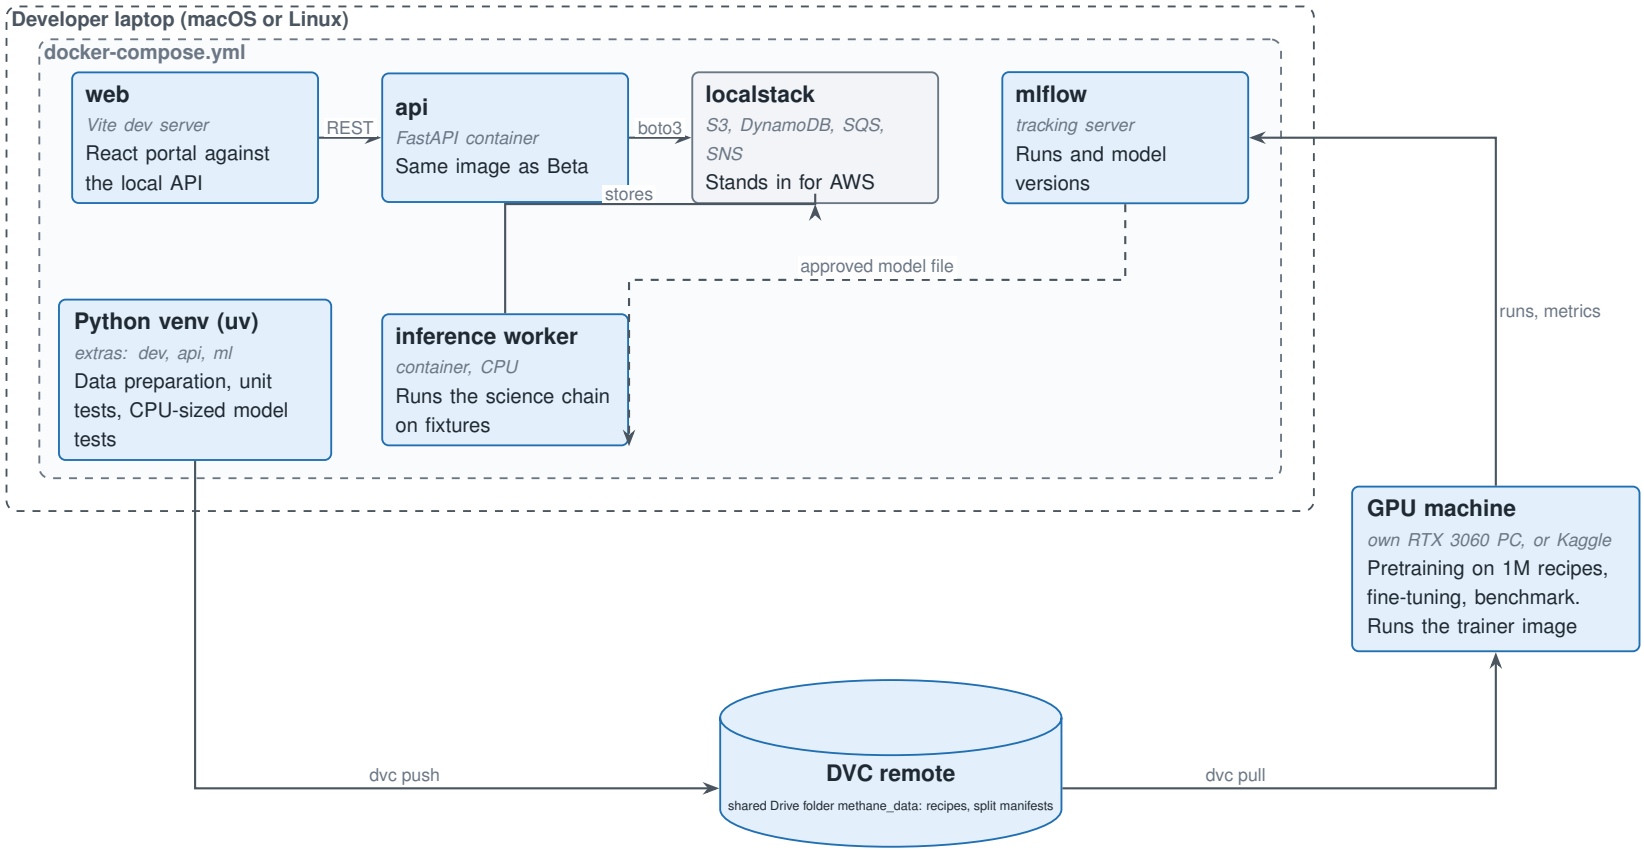

Data preparation stays on the Mac and writes to the shared Drive folder through DVC; the GPU machine pulls it. A model reaches the inference worker only as a registered, approved file, never by copying a notebook output. Sentinel-2 STAC and ERA5 are called only by runs that do not use fixtures.