Eight candidates, one interface: a fair comparison is built into the code

Level 03 - model zoo. All models take the same tensor, return logits at the same resolution, and are trained and scored by the same code on the same split.

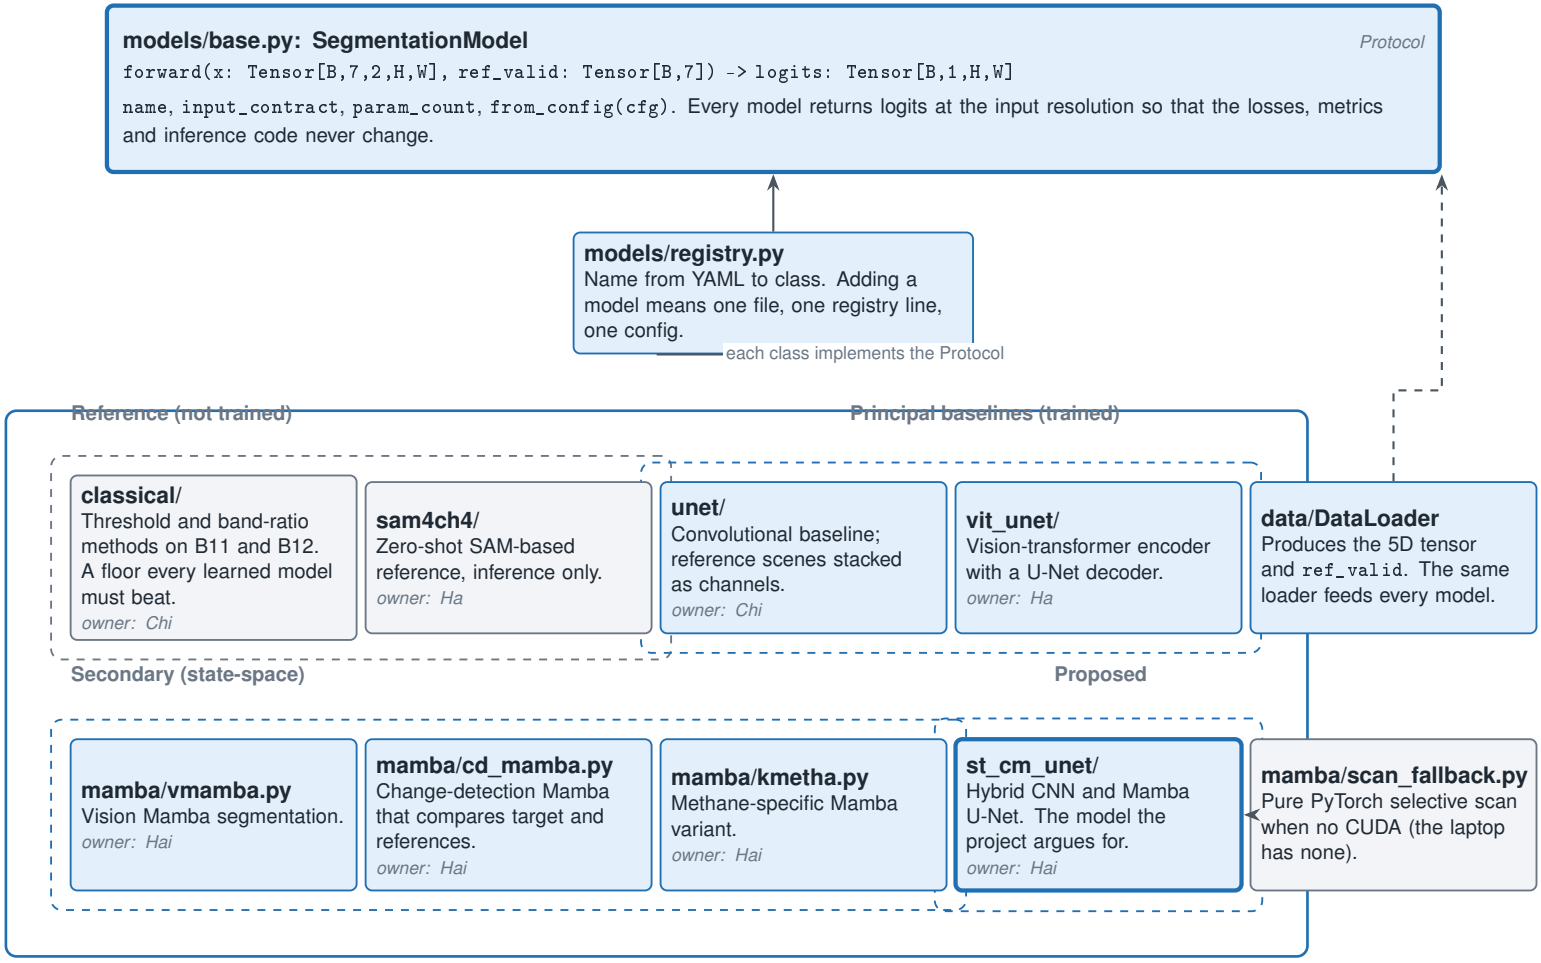

Why this shape: the benchmark compares models, so the comparison itself must not depend on per-model glue. Anything a model needs that others do not (the Mamba scan) lives inside that model's folder. Owners are proposals from the report's team table, confirmed at sprint planning.